

Master UNIR

K-prototypes

```
In [499]: #!pip install umap  
#!pip install umap-learn
```

```
In [3]: #Importación de librerías necesarias  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import statsmodels.api as sm  
import seaborn as sns  
from datetime import datetime as dt  
from datetime import date  
from kmodes.kprototypes import KPrototypes  
from kmodes.util.dissim import matching_dissim  
from kmodes.util.dissim import euclidean_dissim  
from sklearn.preprocessing import StandardScaler  
from sklearn.cluster import KMeans  
import umap.umap_ as umap  
from sklearn.model_selection import cross_val_score  
import logging  
import shap
```

Realizamos como primer paso la importación del dataset a utilizar:

```
In [4]: #Código para cargar el Dataset  
#date=pd.read_csv("Laboratorio_dataset_car.csv", sep=";")  
data=pd.read_csv("202112_1-1.csv")  
dscopy=data
```

```
In [5]: pd.set_option('display.max_columns', 30)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 95151 entries, 0 to 95150
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SK_CLIENTE                           95151 non-null  object
1   NO_PRODUCTO                           77616 non-null  float64
2   SK_TRANSACCION                         77616 non-null  float64
3   DS_NOMBRE                             77616 non-null  object
4   SK_FE_TRANSACCION                     77616 non-null  float64
5   SK_PRODUCTO_SERVICIO                 77616 non-null  float64
6   DS_LINEA_PRODUCTO                    77616 non-null  object
7   DS_FAMILIA                           77616 non-null  object
8   DS_CLASE                             77616 non-null  object
9   VALOR_MVTO                           77616 non-null  float64
10  CANTIDAD_TRANSACCION                  77616 non-null  float64
11  DS_SECTOR_PIB                         77616 non-null  object
12  DS_TIPO_EMPRESA                       77616 non-null  object
13  DS_TIPO_PERSONA                       77616 non-null  object
14  DS_GENERO                             77616 non-null  object
15  FE_VINCULACION_CLIENTE                77606 non-null  object
16  DS_ESTADO_CIVIL                       77616 non-null  object
17  FE_NACIMIENTO                          70247 non-null  object
18  DS_PAIS_NACIMIENTO                    77616 non-null  object
19  DS_NIVEL_ESTUDIOS                     77616 non-null  object
20  DS_PROFESION                           77616 non-null  object
21  DS_OCUPACION                           77616 non-null  object
22  DS_TIPO_VIVIENDA                       77616 non-null  object
23  DS_RAZON_SOCIAL                        77616 non-null  object
24  ID_CLIENTE                            77616 non-null  float64
25  NO_PERSONAS_A_CARGO                   77616 non-null  float64
26  REF_NUM                               58685 non-null  float64
27  FE_APERTURA                           77616 non-null  object
28  DK_PERSONA                             77616 non-null  object
29  SK_RC                                 77616 non-null  float64
dtypes: float64(10), object(20)
memory usage: 21.8+ MB
```

```
In [503]: data.head(5)
```

Out[503]:

	SK_CLIENTE	NO_PRODUCTO	SK_TRANSACCION	DS_NOMBRE	SK_FE_TRANSACCION	SK
0	2582480	1.100141e+11	714.0	IVA SOBRE CHEQUERAS (IGUAL SER	20211206.0	
1	2582480	1.100141e+11	758.0	COBRO DE CHEQUERA (IGUAL SERIE	20211206.0	
2	2582480	1.100141e+11	644.0	ND.TIMBRES CHEQUERA	20211206.0	
3	2582480	1.100141e+11	2562.0	CONSIGNACION AVAL	20211206.0	
4	2582480	1.100141e+11	345.0	ABONOS POR A.C.H	20211206.0	

Procedemos a limpiar las columnas que no contienen información de relevancia para la construcción del modelo y las filas con datos nulos.

```
In [6]: drop_list=["DS_FAMILIA","REF_NUM","ID_CLIENTE","DK_PERSONA","DS_NOMBRE","DS_RAZA",  
                  "DS_TIPO_EMPRESA","SK_RC","DK_PERSONA","DS_PAIS_NACIMIENTO","FE_APER",  
                  "FE_VINCULACION_CLIENTE"]  
data = data.drop(drop_list, axis=1)  
data=data.dropna()  
data=data.drop_duplicates()
```

Además las columnas de fecha son cambiadas por años, con el fin de mejorar la interpretación de éstas en la construcción del modelo.

```
In [7]: data['EDAD'] = (dt.now()-pd.to_datetime(data['FE_NACIMIENTO'],errors='coerce'))  
data['ANTIGUEDAD_CLIENTE'] = (dt.now()-pd.to_datetime(data['FE_VINCULACION_CLIENTE'],errors='coerce'))  
data=data.dropna()
```

Se reordenan las columnas para separarlas en categóricas y numéricas:

```
In [8]: data = data[['SK_CLIENTE',  
                    'SK_PRODUCTO_SERVICIO', 'DS_LINEA_PRODUCTO', 'DS_CLASE',  
                    'CANTIDAD_TRANSACCION', 'DS_GENERO',  
                    'DS_ESTADO_CIVIL',  
                    'DS_NIVEL_ESTUDIOS', 'DS_PROFESION',  
                    'DS_OCUPACION', 'DS_TIPO_VIVIENDA', 'NO_PERSONAS_A_CARGO', 'EDAD', 'ANTIGUO',  
                    'VALOR_MVTO']]
```

Teniendo el dataset limpio, se procede a realizar la curva de codo, con el fin de estimar el valor óptimo para el número de clusters.

```

In [9]: categorical_index = list(range(0,10))
df = data.drop(["SK_CLIENTE"], axis=1).sample(10000, random_state=40)
scaled_X = StandardScaler().fit_transform(df[["NO_PERSONAS_A_CARGO", 'EDAD', 'ANTIGUEDAD_CLIENTE', 'VALOR_MVTO']] = scaled_X)
# Function for plotting elbow curve
def plot_elbow_curve(start, end, data):
    no_of_clusters = list(range(start, end+1))
    cost_values = []

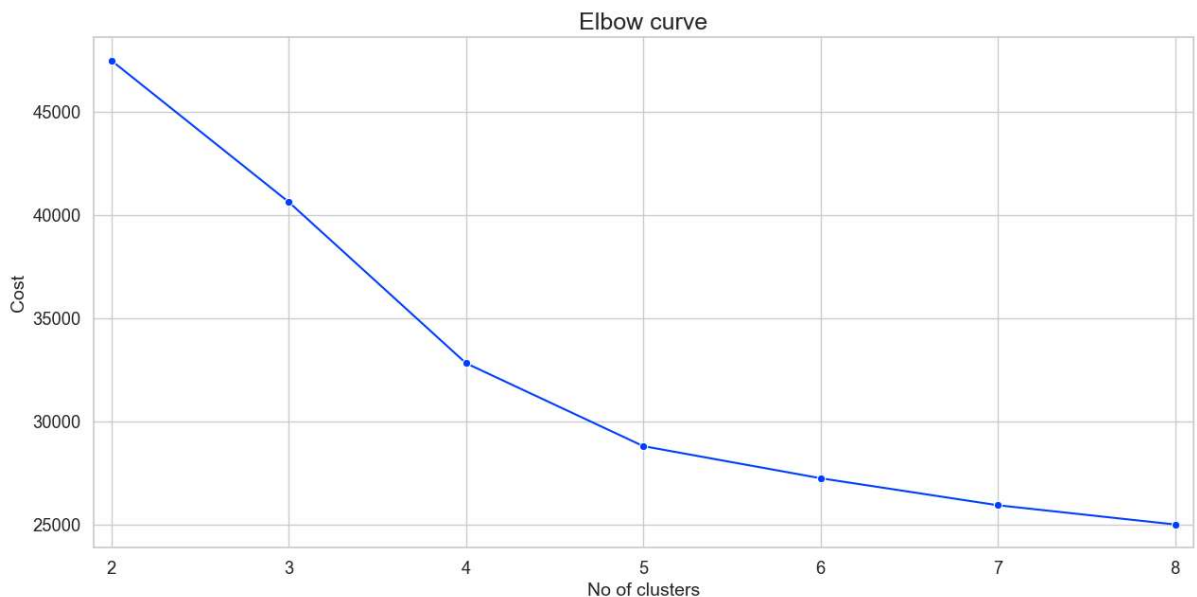
    for k in no_of_clusters:
        test_model = KPrototypes(n_clusters=k, init='Huang', random_state=42)
        test_model.fit_predict(data, categorical=categorical_index)
        cost_values.append(test_model.cost_)

    sns.set_theme(style="whitegrid", palette="bright", font_scale=1.2)

    plt.figure(figsize=(15, 7))
    ax = sns.lineplot(x=no_of_clusters, y=cost_values, marker="o", dashes=False)
    ax.set_title('Elbow curve', fontsize=18)
    ax.set_xlabel('No of clusters', fontsize=14)
    ax.set_ylabel('Cost', fontsize=14)
    ax.set(xlim=(start-0.1, end+0.1))
    plt.plot();

# Plotting elbow curve for k=2 to k=7
plot_elbow_curve(2,8,df)

```



Se utilizará k=4 para la realización del cluster, debido a que es un valor apropiado observando la gráfica. Por tanto, se procede a la construcción del modelo.

```
In [10]: model_4 = KPrototypes(n_clusters=4, init='Huang', random_state=42, n_jobs=-1)
model_4.fit_predict(df, categorical=categorical_index)
labels = model_4.labels_
print(model_4.cost_)
```

32825.607173671764

Con el fin de graficar nuestro dataset, se debe realizar algún método de reducción de dimensionalidad. Se trabajará por medio de reducción con embeddings, utilizando el método UMAP, pues con éste se tiene la particularidad de poder utilizar datos mixtos, de tipo categórico y numérico.

```

In [13]: ##preprocessing categorical
numerical = df[['NO_PERSONAS_A_CARGO', 'EDAD', 'ANTIGUEDAD_CLIENTE', 'VALOR_MVTO']]
categorical = df[['SK_PRODUCTO_SERVICIO', 'DS_LINEA_PRODUCTO', 'DS_CLASE', 'CAN
                'DS_ESTADO_CIVIL', 'DS_NIVEL_ESTUDIOS', 'DS_PROFESION', 'DS_OCUPA
categorical_dummies = pd.get_dummies(categorical)

#Percentage of columns which are categorical is used as weight parameter in emb
categorical_weight = len(categorical) / df.shape[1]

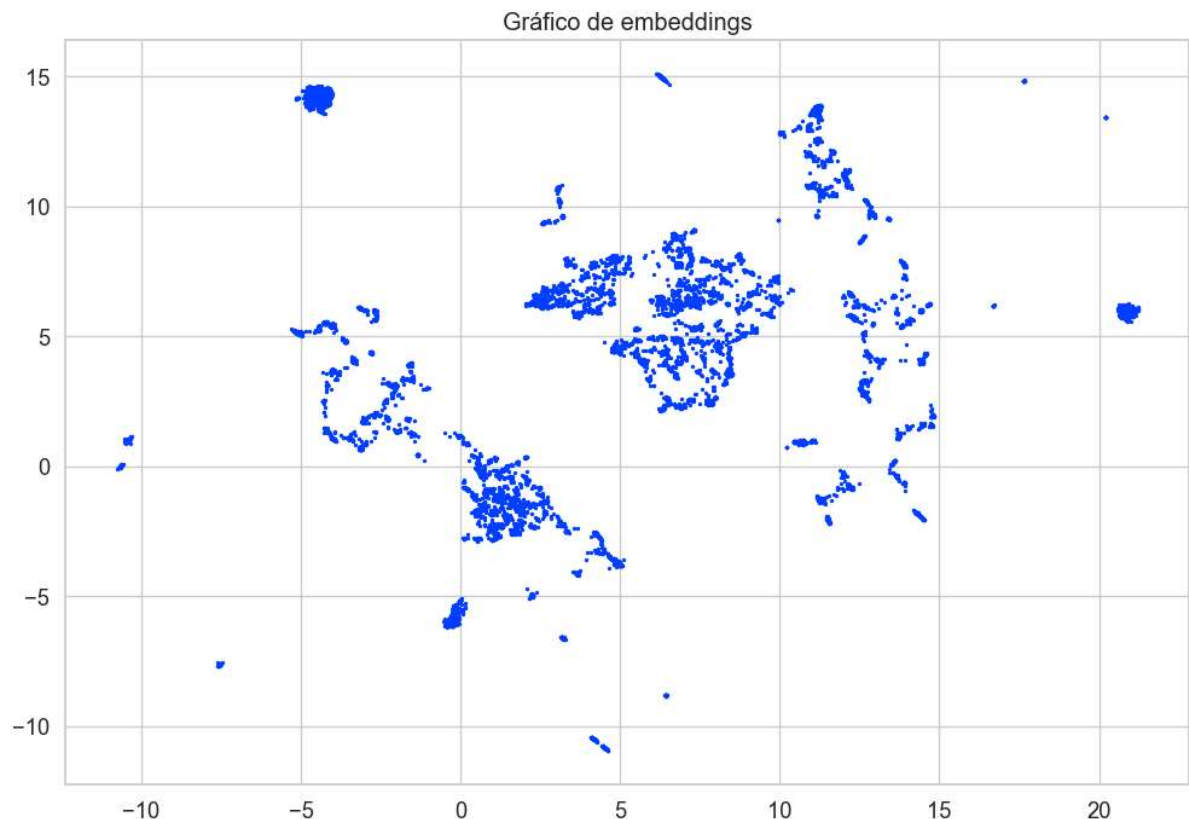
#Embedding numerical & categorical
fit1 = umap.UMAP(metric='l2').fit(numerical)
fit2 = umap.UMAP(metric='dice').fit(categorical_dummies)

#Augmenting the numerical embedding with categorical
intersection = umap.general_simplicial_set_intersection(fit1.graph_, fit2.graph_
intersection = umap.reset_local_connectivity(intersection)
embedding = umap.simplicial_set_embedding(fit1._raw_data, intersection, fit1.n_
                                         fit1._initial_alpha, fit1._a, f
                                         fit1.repulsion_strength, fit1.r
                                         200, 'random', np.random, fit1.
                                         fit1._metric_kwds, False, densn

plt.figure(figsize=(12, 8))
x,y = zip(*embedding[0])
plt.scatter(x,y, s=2, cmap='Spectral', alpha=1.0)
plt.title('Gráfico de embeddings')
plt.show()

```


gradient function is not yet implemented for dice distance metric; inverse_transform will be unavailable
No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored



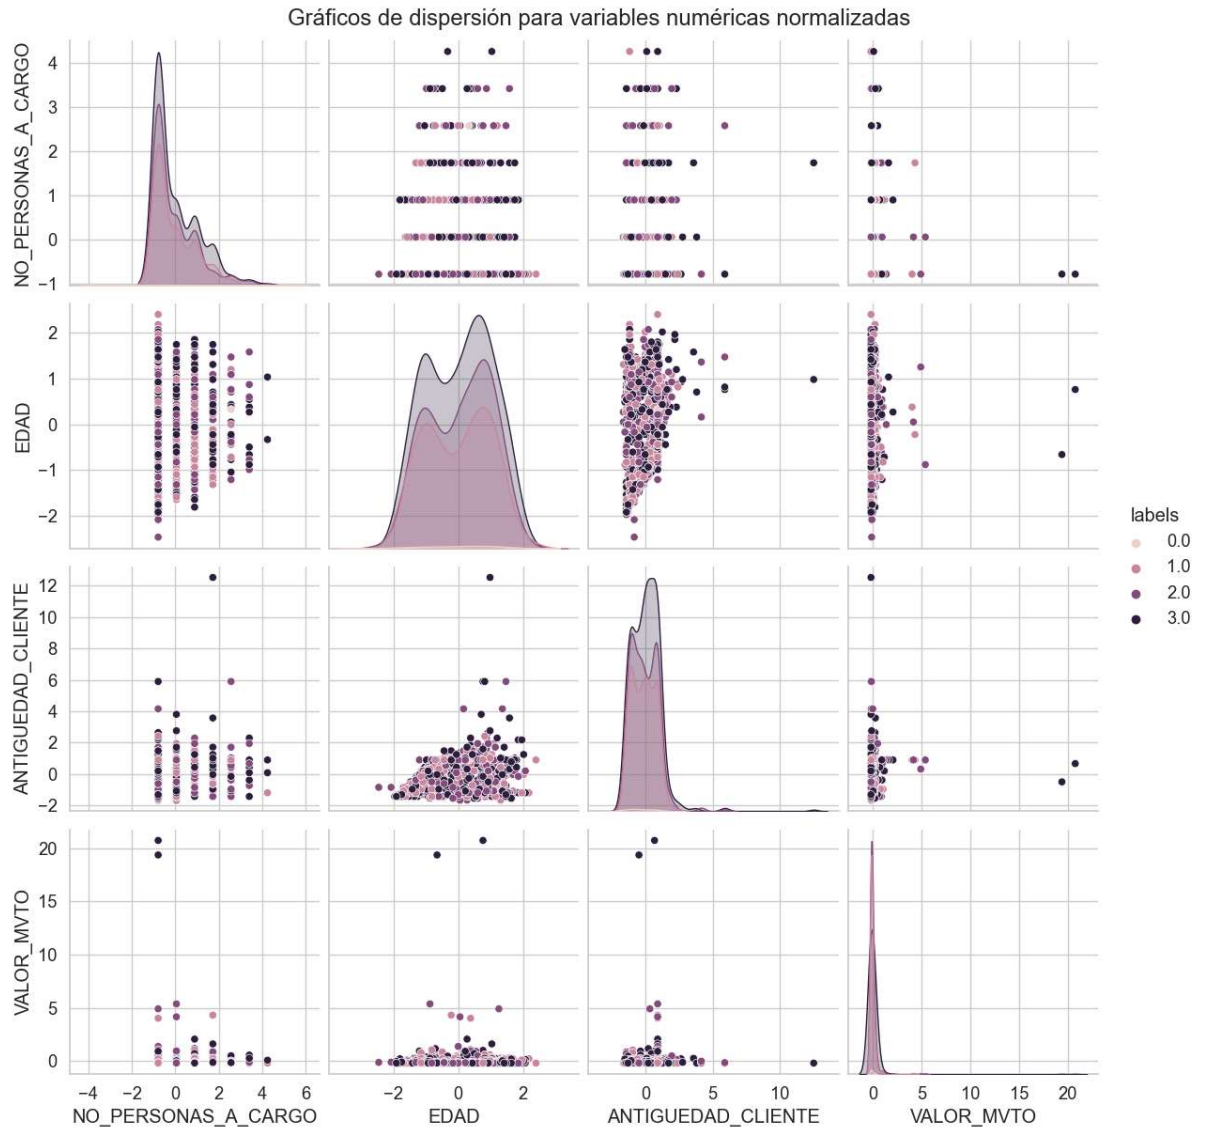
Ahora que ya se tiene el gráfico del embedding, es importante revisar cómo se aprecia la separación de clusters en éste.

```
In [14]: fig, ax = plt.subplots(figsize=(12, 6))
x,y = zip(*embedding[0])
scatter = ax.scatter(x,y,s=2, cmap='Spectral', alpha=1.0, c=labels)
legend1 = ax.legend(*scatter.legend_elements(),
                    loc="upper right", title="Classes")
ax.add_artist(legend1)
ax.set_title('Gráfico de embeddings con clusters',fontsize=15)
```

Out[14]: Text(0.5, 1.0, 'Gráfico de embeddings con clusters')




```
In [15]: df_pairplot = numerical.copy()
df_pairplot.loc[:, "labels"] = pd.Series(labels)
ax = sns.pairplot(df_pairplot, hue='labels', height=3, aspect=1)
ax.fig.suptitle('Gráficos de dispersión para variables numéricas normalizadas',
sns.set_context("paper", rc={"axes.labelsize":10})
```



Ahora que tenemos los clusters definidos, realicemos una validación de estos por medio del coeficiente de silueta y tener una medida de qué tan apropiado resultó el agrupamiento.

```
In [19]: df_model4 = pd.concat([categorical_dummies, numerical], axis=1, ignore_index=1)
from sklearn.metrics import silhouette_score
silhouette_score(df_model4, labels)
```

Out[19]: 0.15703955461258348

```
In [20]: from scipy.stats import chi2_contingency

def cramers_V(var1, var2):
    crosstab = np.array(pd.crosstab(var1, var2))
    stats = chi2_contingency(crosstab)[0]
    cram_V = stats / (np.sum(crosstab) * (min(crosstab.shape) - 1))
    return cram_V

def cramers_col(column_name):
    col = pd.Series(np.empty(df.columns.shape), index=df.columns, name=column_name)
    for row in df:
        cram = cramers_V(df[column_name], df[row])
        col[row] = round(cram, 2)
    return col

df.apply(lambda column: cramers_col(column.name))
```

Out[20]:

	SK_PRODUCTO_SERVICIO	DS_LINEA_PRODUCTO	DS_CLASE	CANTIDAD_TRANSACCION
SK_PRODUCTO_SERVICIO	1.00	1.00	1.00	
DS_LINEA_PRODUCTO	1.00	1.00	1.00	
DS_CLASE	1.00	1.00	0.99	
CANTIDAD_TRANSACCION	0.00	0.00	0.00	
DS_GENERO	0.00	0.00	0.00	
DS_ESTADO_CIVIL	0.00	0.00	0.00	
DS_NIVEL_ESTUDIOS	0.00	0.00	0.00	
DS_PROFESION	0.02	0.02	0.06	
DS_OCUPACION	0.03	0.03	0.11	
DS_TIPO_VIVIENDA	0.00	0.00	0.00	
NO PERSONAS A CARGO	0.00	0.00	0.00	

Otra forma de evaluar internamente qué tan eficiente es el cluster, es utilizando el coeficiente F1 de validación cruzada, para esto nos apoyaremos en un clasificador LGBM, que nos permite trabajar variables de tipo mixto.


```
In [22]: from lightgbm import LGBMClassifier
for c in df.select_dtypes(include='object'):
    df[c] = df[c].astype('category')
clf_kp = LGBMClassifier(colsample_by_tree=0.8)
cv_scores_kp = cross_val_score(clf_kp, df, labels, scoring='f1_weighted')
print(f'CV F1 score for K-Prototypes clusters is {np.mean(cv_scores_kp)}')
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

Dado que el valor de validación cruzada es 0.985, se puede establecer que el dataset fue clasificado en grupos con una separación distinguible y representativa. Podemos ahora revisar la significancia de cada una de las variables usadas en el clustering, para lo cual primero entrenaremos un clasificador y luego apoyandonos en los valores SHAP, se revisará la importancia de cada variable en el cluster.

```
In [23]: clf_kp.fit(df, labels)
explainer_kp = shap.TreeExplainer(clf_kp)
shap_values_kp = explainer_kp.shap_values(df)
shap.summary_plot(shap_values_kp, df, plot_type="bar", plot_size=(15, 10))
```

[illegible]



Basándonos en la gráfica realizada, repetiremos el cluster usando solamente aquellas variables que demuestran significancia.

```

In [25]: categorical_index = list(range(0,5))
drop_list2= ["CANTIDAD_TRANSACCION", "DS_PROFESION", "SK_PRODUCTO_SERVICIO", "DS_C
df2 =df.drop(drop_list2, axis=1)

# Function for plotting elbow curve
def plot_elbow_curve(start, end, data):
    no_of_clusters = list(range(start, end+1))
    cost_values = []


    for k in no_of_clusters:
        test_model = KPrototypes(n_clusters=k, init='Huang', random_state=42)
        test_model.fit_predict(data, categorical=categorical_index)
        cost_values.append(test_model.cost_)

    sns.set_theme(style="whitegrid", palette="bright", font_scale=1.2)

    plt.figure(figsize=(15, 7))
    ax = sns.lineplot(x=no_of_clusters, y=cost_values, marker="o", dashes=False)
    ax.set_title('Elbow curve', fontsize=18)
    ax.set_xlabel('No of clusters', fontsize=14)
    ax.set_ylabel('Cost', fontsize=14)
    ax.set(xlim=(start-0.1, end+0.1))
    plt.plot();

# Plotting elbow curve for k=2 to k=7
plot_elbow_curve(2,8,df2)

```



```

In [26]: model_val = KPrototypes(n_clusters=4, init='Huang', random_state=42, n_jobs=-1)
model_val.fit_predict(df2, categorical=categorical_index)
labels_2 = model_val.labels_
print(model_val.cost_)

```

17813.23543781217

```

In [29]: ##preprocessing categorical
numerical_2 = df2[['EDAD', 'ANTIGUEDAD_CLIENTE', 'VALOR_MVTO']]
categorical_2 = df2[['DS_LINEA_PRODUCTO', 'DS_GENERO',
                    'DS_ESTADO_CIVIL', 'DS_OCUPACION', 'DS_TIPO_VIVIENDA']]
categorical_dummies_2 = pd.get_dummies(categorical_2)

#Percentage of columns which are categorical is used as weight parameter in emb
categorical_weight = len(categorical_2) / df2.shape[1]

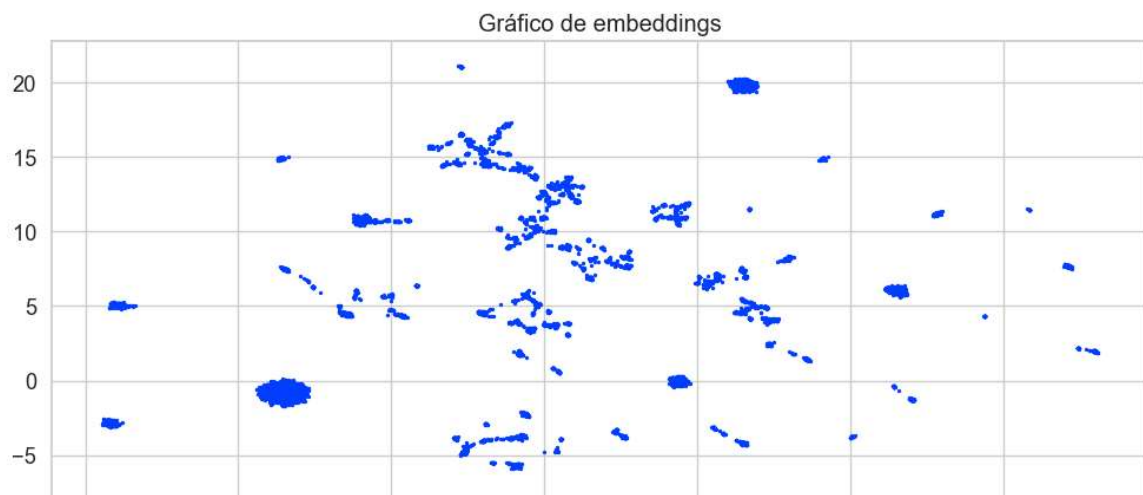
#Embedding numerical & categorical
fit1 = umap.UMAP(metric='l2').fit(numerical_2)
fit2 = umap.UMAP(metric='dice').fit(categorical_dummies_2)

#Augmenting the numerical embedding with categorical
intersection = umap.general_simplicial_set_intersection(fit1.graph_, fit2.graph_)
intersection = umap.reset_local_connectivity(intersection)
embedding = umap.simplicial_set_embedding(fit1._raw_data, intersection, fit1.n_
                                         fit1._initial_alpha, fit1._a, f
                                         fit1.repulsion_strength, fit1.r
                                         200, 'random', np.random, fit1.
                                         fit1._metric_kwds, False, densn

plt.figure(figsize=(12, 6))
x,y = zip(*embedding[0])
plt.scatter(x,y, s=2, cmap='Spectral', alpha=1.0)
plt.title('Gráfico de embeddings')
plt.show()


```

gradient function is not yet implemented for dice distance metric; inverse_transform will be unavailable
No data for colormapping provided via 'c'. Parameters 'cmap' will be ignored




```
In [30]: fig, ax = plt.subplots(figsize=(12, 6))
x,y = zip(*embedding[0])
scatter = ax.scatter(x,y,s=2, cmap='Spectral', alpha=1.0, c=labels_2)
legend1 = ax.legend(*scatter.legend_elements(),
                    loc="upper right", title="Classes")
ax.add_artist(legend1)
ax.set_title('Gráfico de embeddings con clusters', fontsize=15)
```

Out[30]: Text(0.5, 1.0, 'Gráfico de embeddings con clusters')



```
In [31]: df_model_val = pd.concat([categorical_dummies_2, numerical_2], axis=1, ignore_index=True)
from sklearn.metrics import silhouette_score
silhouette_score(df_model_val, labels_2)
```

Out[31]: 0.1980837034323855

Como se puede apreciar el valor del coeficiente de silueta tiene una mejora considerable, demostrando una optimización del cluster.

```
In [33]: from lightgbm import LGBMClassifier
for c in df2.select_dtypes(include='object'):
    df2[c] = df[c].astype('category')
clf_kp = LGBMClassifier(colsample_by_tree=0.8)
cv_scores_kp = cross_val_score(clf_kp, df2, labels_2, scoring='f1_weighted')
print(f'CV F1 score for K-Prototypes clusters is {np.mean(cv_scores_kp)}')
```

```
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
```

El coeficiente de validación cruzada se mantiene en un valor parecido, por lo cual internamente con respecto al dataset el cluster hace una segmentación adecuada.